

**Stat 5810, Section 003**  
**Statistical Visualization I**  
**Fall 2018**

**Dr. Jürgen Symanzik**

Utah State University

Department of Mathematics and Statistics

3900 Old Main Hill

Logan, UT 84322-3900

Tel.: (435) 797-0696

FAX: (435) 797-1822

e-mail: [symanzik@math.usu.edu](mailto:symanzik@math.usu.edu)

Web: <http://www.math.usu.edu/~symanzik/>



# Contents

## Acknowledgements

This course uses some of the course materials provided by Dr. Mike Minnotte (formerly USU, now with the University of North Dakota) as held in the Fall 2006 semester. Additional materials have been taken from other Statistical Graphics courses, such as the ones offered by Dr. Di Cook (formerly Iowa State University; now Monash University: <http://dicook.org/>) and Dr. Dan Carr (George Mason University: <http://mason.gmu.edu/~dcarr/>). Other examples and R code originate from Heike Hofmann, Paul Murrell, Carson Sievert, Martin Theus, Antony Unwin, Simon Urbanek, Hadley Wickham, Lee Wilkinson, and others. We are likely to include parts from additional authors and sources that will be specified later during the semester.

Thanks are also due to 60+ students and guests who took the former “Stat 6560: Graphical Methods” and the current “Statistical Visualization I & II” courses with me since the Spring 2009 semester for their valuable comments that helped to improve, correct, and extend these lecture notes.

Jürgen Symanzik, September 10, 2018.

# 1 Univariate Plots

## 1.1 Histograms

Example 1:

Four histograms of the same data set, showing the weights in pounds of 132 professional male athletes.

Question:

What can we conclude about the underlying data? And which of these four histograms best reveals this fact?

### Example 2:

An interactive applet that allowed to change the number of classes in a histogram via a slider could be found at <http://www.stat.sc.edu/~west/javahtml/Histogram.html>:

#### **“Histogram Applet:**

This applet is designed to teach students how bin widths (or the number of bins) affect a histogram. The histogram below is for the Old Faithful data set. The observations are the duration (in minutes) for eruptions of the Old Faithful geyser in Yellowstone National Park. Students should interactively change the bin width by dragging the arrow underneath the bin width scale. For large bin widths, the bimodal nature of the dataset is hidden, and for small bin widths the plot reduces to a spike at each data point. What bin width do you think provides the best picture of the underlying data?”

While the URL above no longer exists, the following applet from the Rossman and Chance Applet Collection has a similar functionality: <http://www.rossmanchance.com/applets/Dotplot.html>

First prepare the *faithful* data in R so it appears in a single column without any additional spaces before/after the values:

```
> for (i in 1:length(faithful$eruptions)) {  
+   cat(paste0(faithful$eruptions[i], "\n"))  
+ }
```

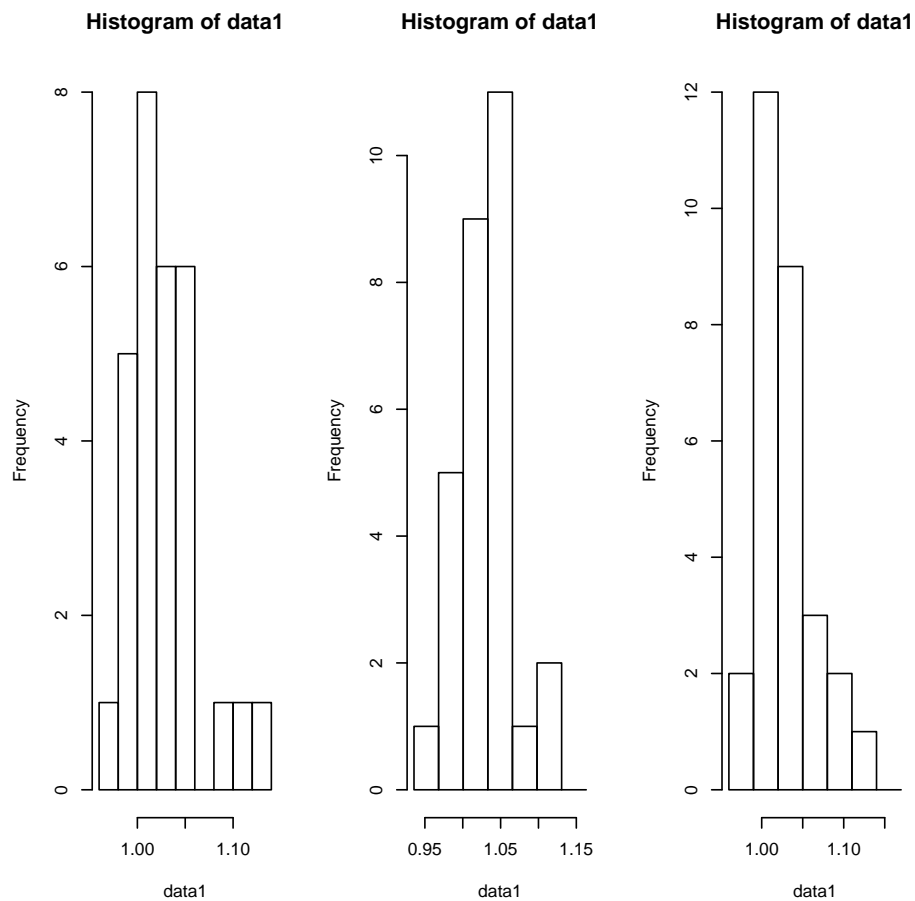
Then copy the data from your R console and paste them into the “Sample data” box on the web site. Click on “Use data” and select “Histogram”. Use the slider and change the “Number of bins”. What happens when this number increases?

### Example 3:

Example data sets from ?:

```
> # Weber (2008), Set 1:
>
> data1 <- c(.968, .982, .991, .993, .998, .999, 1.004, 1.004,
+ 1.007, 1.010, 1.012, 1.015, 1.017, 1.019, 1.021,
+ 1.035, 1.037, 1.037, 1.039, 1.039, 1.042, 1.042,
+ 1.047, 1.053, 1.055, 1.059, 1.081, 1.107, 1.1305)
> par(mfrow = c(1, 3))
> hist(data1) #Default
> hist(data1, breaks = 0.9356 + (0:7) * 0.0325) #A
> hist(data1, breaks = 0.9600 + (0:7) * 0.0300) #B
> summary(data1)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.968	1.004	1.021	1.029	1.042	1.131




```
> # Weber (2008), Set 2:
>
> data2 <- c(2.05, 2.27, 2.50, 2.95, 3.18, 3.41, 3.64, 3.86, 4.09, 4.32,
+ 5.68, 5.91, 6.14, 6.36, 6.59, 6.82, 7.05, 7.50, 7.73, 7.95)
> par(mfrow = c(2, 6))
> hist(data2)
> hist(data2, breaks = 1.425 + (0:3) * 3.2075) #C
> hist(data2, breaks = -1.048 + (0:3) * 3.2075) #D
> hist(data2, breaks = 1.9767 + (0:4) * 1.9789) #E
> hist(data2, breaks = 0.1078 + (0:4) * 1.9789) #F
> hist(data2, breaks = 1.9829 + (0:5) * 1.4750) #G
> hist(data2, breaks = 0.6421 + (0:5) * 1.4750) #H
> hist(data2, breaks = 1.9619 + (0:4) * 1.9060) #I
> hist(data2, breaks = 0.8944 + (0:4) * 1.9060) #J
> hist(data2, breaks = -0.6800 + (0:4) * 2.8400) #K
```




```
> hist(data2, breaks = 1.9542 + (0:4) * 1.5229) #L
> hist(data2, breaks = 1.9619 + (0:7) * .9060) #New 1
> summary(data2)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
2.050	3.353	5.000	5.000	6.647	7.950

**Histogram of dHistogram of dHistogram of dHistogram of dHistogram of dHistogram of d:**



**Histogram of dHistogram of dHistogram of dHistogram of dHistogram of dHistogram of d:**



**Conclusion:** “Never believe any statistics you haven’t falsified yourself.”

(<http://sanmateorealestateblog.com/real-estate/statistics-real-estate/never-believe-any-statistics-you-havent-falsified-yourself/>)

The R help page for hist indicates:

“The generic function hist computes a histogram of the given data values.”

The R help page for the Iris data set indicates:

“This famous (Fisher’s or Anderson’s) iris data set gives the measurements in centimeters of the variables sepal length and width and petal length and width, respectively, for 50 flowers from each of 3 species of iris. The species are *Iris setosa*, *versicolor*, and *virginica*.

iris is a data frame with 150 cases (rows) and 5 variables (columns) named Sepal.Length, Sepal.Width, Petal.Length, Petal.Width, and Species.”

Choosing the number of classes for a histogram:

As seen in the previous three examples, a bad choice for the number of classes (nclass or breaks in the R command) in a histogram or the starting point of an interval and its width can almost entirely hide the most interesting information of the underlying data.

Several suggestions for the number of classes exist and are summarized in ?, p. 112. We define  $\text{range} = x_{(n)} - x_{(1)}$ , where  $n$  represents the number of observations.

- Sturges’ formula (default in R):

$$\text{nclass} = \lceil \log_2 n + 1 \rceil, \quad \text{bin width} = \frac{\text{range}}{\text{nclass}},$$

where  $\lceil \dots \rceil$  indicates the ceiling function.

- Scott’s 1979 formula (“scott” in R):

$$\text{bin width} = 3.5 \hat{\sigma} n^{-1/3}, \quad \text{nclass} = \frac{\text{range}}{\text{bin width}},$$

where  $\hat{\sigma}$  is the estimated standard deviation.

- Freedman and Diaconis 1981 formula (“fd” in R):

$$\text{bin width} = 2 \text{ IQR } n^{-1/3}, \quad \text{nclass} = \frac{\text{range}}{\text{bin width}},$$

where IQR is the inter-quartile range.

- Sometimes, the use of  $n_{class} \approx \sqrt{n}$  is suggested:

<http://www.qimacros.com/qiwizard/how-to-determine-histogram-bin-interval.htm> suggests: *“Take the square root of the number of data points and round up to determine the number of bins required.”*

<http://www.moresteam.com/toolbox/t417.cfm> suggests: *“Calculate the square root of the number of data points and round to the nearest whole number. In the case of our height example, the square root of 50 is 7.07, or 7 when rounded.”*

[http://www.micquality.com/introductory\\_statistics/int08.htm](http://www.micquality.com/introductory_statistics/int08.htm) states: *“There are various ways of calculating the number of bins. I find that using the square root of the number of data values gives as good a result as the more complicated methods. The value is usually on the low side, but you can adjust it upwards to get convenient bin boundaries. Treat the calculated number of bins as a starting point, and adjust it as necessary to give the result you prefer.”*

### Example:

```
> head(iris)
```

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
1	5.1	3.5	1.4	0.2	setosa
2	4.9	3.0	1.4	0.2	setosa
3	4.7	3.2	1.3	0.2	setosa
4	4.6	3.1	1.5	0.2	setosa
5	5.0	3.6	1.4	0.2	setosa
6	5.4	3.9	1.7	0.4	setosa

```
> plength <- iris[, 3]
> n <- length(plength)
> par(mfrow = c(2, 2))
> hist(plength, freq = FALSE,
+      main = "Default (Sturges) Breaks")
> hist(plength, breaks = as.integer(sqrt(n)), freq = FALSE,
+      main = "sqrt(n) Breaks")
> hist(plength, breaks = "scott", freq = FALSE,
+      main = "Scott Breaks")
> hist(plength, breaks = "fd", freq = FALSE,
```

```
+      main = "FD Breaks")
```

```
> nclass.Sturges(plength)
```

```
[1] 9
```

```
> sqrt(n)
```

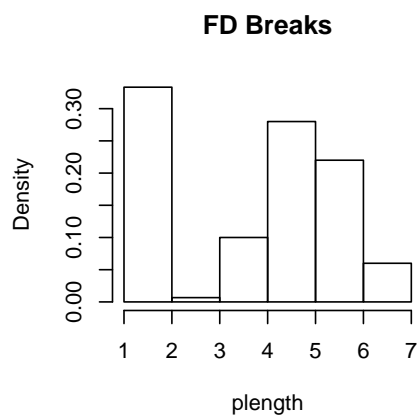


```
[1] 12.24745
```

```
> nclass.scott(plength)
```

```
[1] 6
```

```
> nclass.FD(plength)
```

```
[1] 5
```

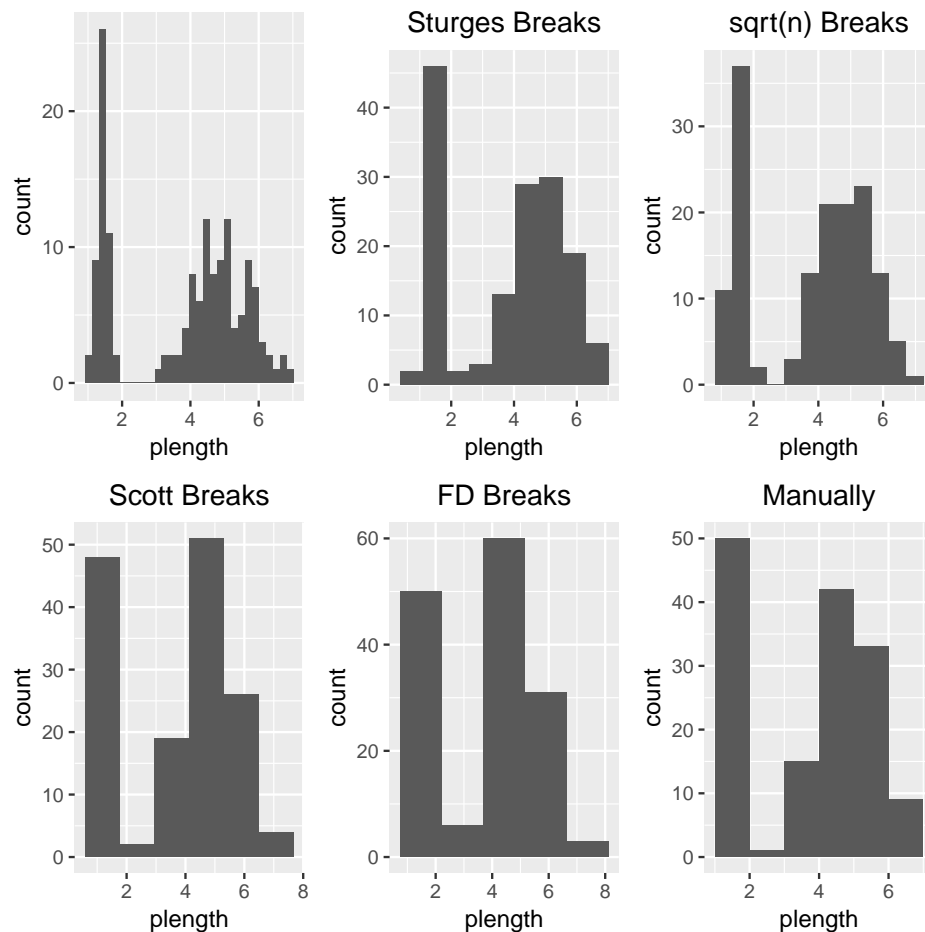


Now in ggplot2:

```

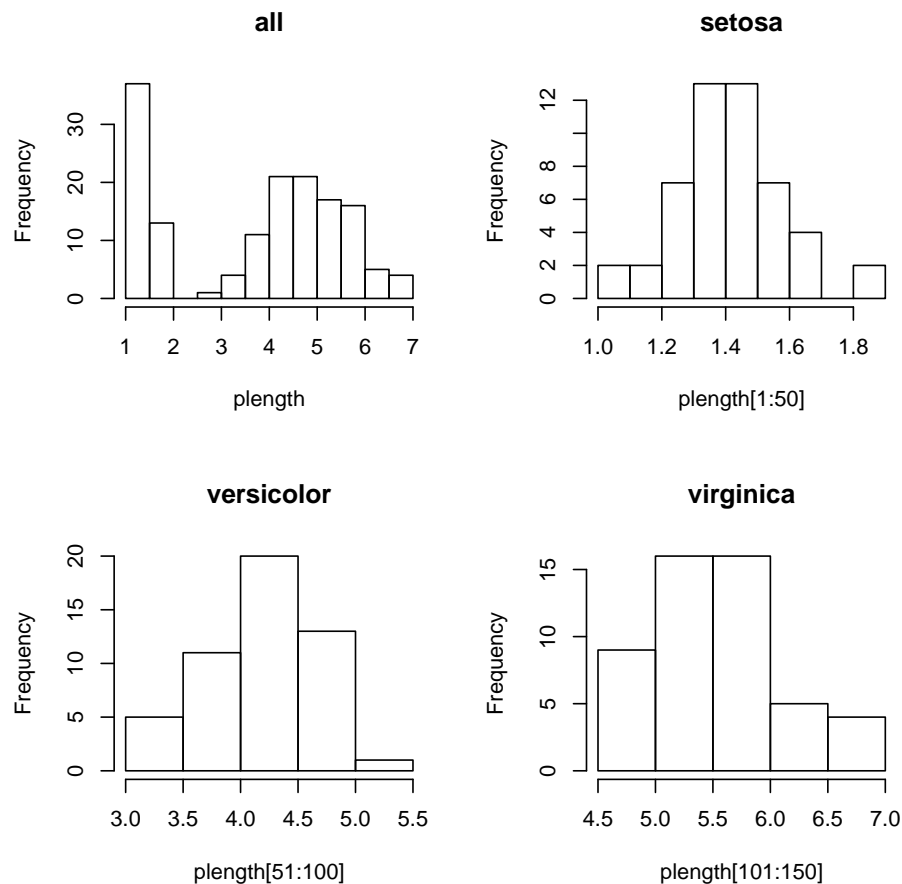
> library(ggplot2)
> library(gridExtra)
> h1 <- ggplot(iris, aes(x = plength)) +
+   geom_histogram()
> h2 <- ggplot(iris, aes(x = plength)) +
+   geom_histogram(bins = nclass.Sturges(plength)) +
+   ggtitle("Sturges Breaks") +
+   theme(plot.title = element_text(hjust = 0.5))
> h3 <- ggplot(iris, aes(x = plength)) +
+   geom_histogram(bins = as.integer(sqrt(n))) +
+   ggtitle("sqrt(n) Breaks") +
+   theme(plot.title = element_text(hjust = 0.5))
> h4 <- ggplot(iris, aes(x = plength)) +
+   geom_histogram(bins = nclass.scott(plength)) +
+   ggtitle("Scott Breaks") +
+   theme(plot.title = element_text(hjust = 0.5))
> h5 <- ggplot(iris, aes(x = plength)) +
+   geom_histogram(bins = nclass.FD(plength)) +
+   ggtitle("FD Breaks") +
+   theme(plot.title = element_text(hjust = 0.5))
> h6 <- ggplot(iris, aes(x = plength)) +
+   geom_histogram(binwidth = 1.0, center = 0.5) +
+   ggtitle("Manually") +
+   theme(plot.title = element_text(hjust = 0.5))
> grid.arrange(h1, h2, h3, h4, h5, h6, nrow = 2)

```



Finally, how do the histograms for the three species look like?

```
> par(mfrow = c(2, 2))
> hist(length, main = "all")
> hist(length[1:50], main = "setosa")
> hist(length[51:100], main = "versicolor")
> hist(length[101:150], main = "virginica")
```



**Note:**

- These various methods (Sturges, Scott, FD, sqrt) provide suggestions for the number of classes only. To enforce particular breaks, we have to provide a vector giving the exact break points between the histogram cells. However, good software will use the suggestions and then make further adjustments to obtain meaningful class breaks for a human reader, e.g., use integers (and multiples of 5 or 10, etc.) as the boundaries.
- Carefully check whether class intervals are left-open or right-open. R class intervals by default are left-open whereas most readers prefer right-open intervals. Also, check in which class interval the minimum and maximum of a data set are included. For continuous data, there will be little differences in the appearance of a histogram, but for discrete data, different settings may result in a dramatically different visual appearance of a histogram. R provides arguments (`include.lowest` and `right`) to adjust these options.

- ?, p. 30, state: “Since it is relatively complicated both to draw and to read histograms with classes of different size we recommend that, as far as possible, both tables and charts should be made with classes of equal length.”



## 1.2 Averaged Shifted Histograms

?, p. 307, states:

“? provides a general overview on techniques for density estimation, including averaged shifted histograms (ASH) and kernel density estimators, including possible visualization techniques via contour surfaces, (transparent)  $\alpha$ -level contours, and contour shells.”

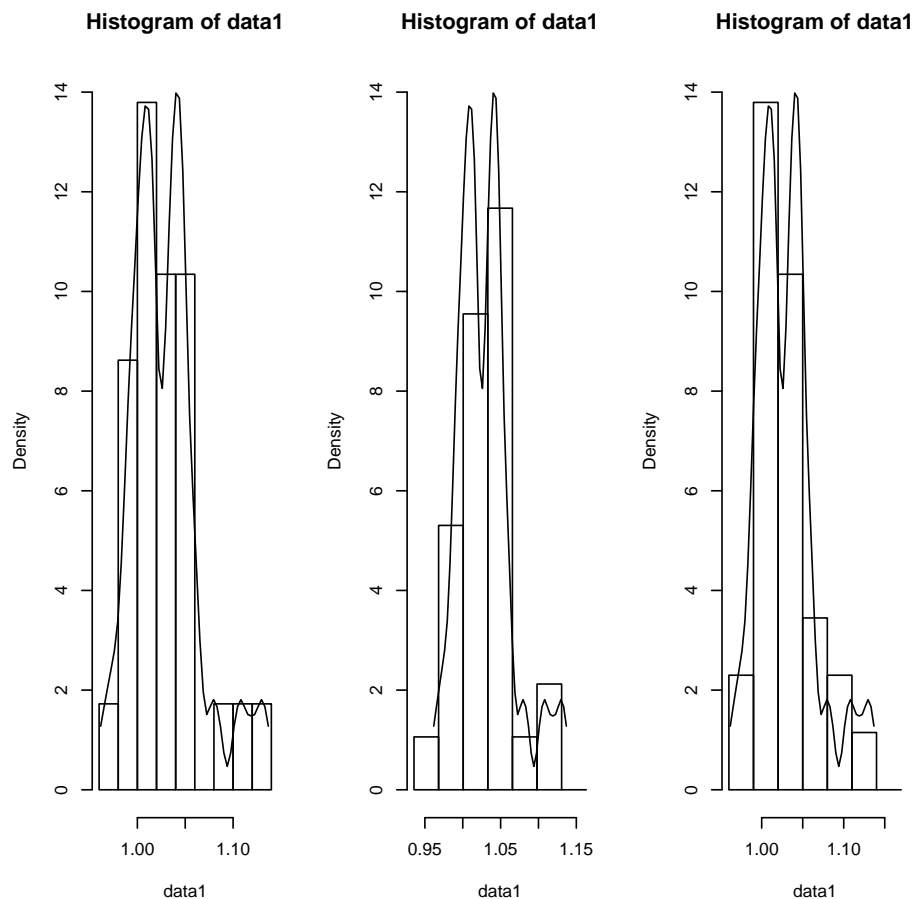
ASH plots were originally introduced in ?. They are created by averaging several shifted histograms and further smoothing the result. Details are provided in Chapter 5 of ?.

ASH plots may not be easy to explain to non-statisticians, but they may help to determine which histograms may be closest to the underlying data.

```
> # Weber (2008), Set 1:
>
> data1 <- c(.968, .982, .991, .993, .998, .999, 1.004, 1.004,
+ 1.007, 1.010, 1.012, 1.015, 1.017, 1.019, 1.021,
+ 1.035, 1.037, 1.037, 1.039, 1.039, 1.042, 1.042,
+ 1.047, 1.053, 1.055, 1.059, 1.081, 1.107, 1.1305)
> library(ash)
> f1 <- ash1(bin1(data1, nbin = 50), 5) # compute ash estimate

[1] "ash estimate nonzero outside interval ab"

> par(mfrow = c(1, 3))
> hist(data1, freq = FALSE, ylim = c(0, 14)) #Default
> lines(f1, type = "l") # line plot of estimate
> hist(data1, breaks = 0.9356 + (0:7) * 0.0325, freq = FALSE, ylim = c(0, 14)) #A
> lines(f1, type = "l") # line plot of estimate
> hist(data1, breaks = 0.9600 + (0:7) * 0.0300, freq = FALSE, ylim = c(0, 14)) #B
> lines(f1, type = "l") # line plot of estimate
```



```
> # Weber (2008), Set 2:
```

```
>
```

```
> data2 <- c(2.05, 2.27, 2.50, 2.95, 3.18, 3.41, 3.64, 3.86, 4.09, 4.32,  
+ 5.68, 5.91, 6.14, 6.36, 6.59, 6.82, 7.05, 7.50, 7.73, 7.95)
```

```
> f2 <- ash1(bin1(data2, nbin = 50), 5) # compute ash estimate
```

```
[1] "ash estimate nonzero outside interval ab"
```

```
> par(mfrow = c(2, 6))
```

```
> hist(data2, freq = FALSE, ylim = c(0, 0.25))
```

```
> lines(f2, type = "l") # line plot of estimate
```

```
> hist(data2, breaks = 1.425 + (0:3) * 3.2075, freq = FALSE, ylim = c(0, 0.25)) #C
```

```
> lines(f2, type = "l") # line plot of estimate
```

```
> hist(data2, breaks = -1.048 + (0:3) * 3.2075, freq = FALSE, ylim = c(0, 0.25)) #D
```


```
> lines(f2, type = "l") # line plot of estimate
```

```


> hist(data2, breaks = 1.9767 + (0:4) * 1.9789, freq = FALSE, ylim = c(0, 0.25)) #E
> lines(f2 , type = "l") # line plot of estimate
> hist(data2, breaks = 0.1078 + (0:4) * 1.9789, freq = FALSE, ylim = c(0, 0.25)) #F
> lines(f2 , type = "l") # line plot of estimate
> hist(data2, breaks = 1.9829 + (0:5) * 1.4750, freq = FALSE, ylim = c(0, 0.25)) #G
> lines(f2 , type = "l") # line plot of estimate
> hist(data2, breaks = 0.6421 + (0:5) * 1.4750, freq = FALSE, ylim = c(0, 0.25)) #H
> lines(f2 , type = "l") # line plot of estimate
> hist(data2, breaks = 1.9619 + (0:4) * 1.9060, freq = FALSE, ylim = c(0, 0.25)) #I
> lines(f2 , type = "l") # line plot of estimate
> hist(data2, breaks = 0.8944 + (0:4) * 1.9060, freq = FALSE, ylim = c(0, 0.25)) #J
> lines(f2 , type = "l") # line plot of estimate
> hist(data2, breaks = -0.6800 + (0:4) * 2.8400, freq = FALSE, ylim = c(0, 0.25)) #K
> lines(f2 , type = "l") # line plot of estimate
> hist(data2, breaks = 1.9542 + (0:4) * 1.5229, freq = FALSE, ylim = c(0, 0.25)) #L
> lines(f2 , type = "l") # line plot of estimate
> hist(data2, breaks = 1.9619 + (0:7) * .9060, freq = FALSE, ylim = c(0, 0.25)) #New 1
> lines(f2 , type = "l") # line plot of estimate

```

Histogram of dHistogram of dHistogram of dHistogram of dHistogram of dHistogram of d



Histogram of dHistogram of dHistogram of dHistogram of dHistogram of dHistogram of d



### 1.3 Stem-and-Leaf Plots

The R help page for `stem` indicates:

“`stem` produces a stem-and-leaf plot of the values in `x`.”

?, p. 113, further specify: “A *stem-and-leaf plot* is an enhanced histogram. The data are divided into bins, but the ‘height’ is replaced by the next digits in order.”

While stem-and-leaf plots show more details than histograms, they are rarely used in practice, in particular not for larger data sets. However, they are a useful first step to the creation of histograms and therefore appear in many introductory statistics textbooks, e.g., ?, pp. 9–12, and ?, pp. 46–49.

```
> sort(plength)
```

```
[1] 1.0 1.1 1.2 1.2 1.3 1.3 1.3 1.3 1.3 1.3 1.3 1.4 1.4 1.4 1.4 1.4 1.4 1.4
[19] 1.4 1.4 1.4 1.4 1.4 1.4 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5
[37] 1.5 1.6 1.6 1.6 1.6 1.6 1.6 1.6 1.7 1.7 1.7 1.7 1.9 1.9 3.0 3.3 3.3 3.5
[55] 3.5 3.6 3.7 3.8 3.9 3.9 3.9 4.0 4.0 4.0 4.0 4.0 4.1 4.1 4.1 4.2 4.2 4.2
[73] 4.2 4.3 4.3 4.4 4.4 4.4 4.4 4.5 4.5 4.5 4.5 4.5 4.5 4.5 4.5 4.6 4.6 4.6
[91] 4.7 4.7 4.7 4.7 4.7 4.8 4.8 4.8 4.8 4.9 4.9 4.9 4.9 4.9 5.0 5.0 5.0 5.0
[109] 5.1 5.1 5.1 5.1 5.1 5.1 5.1 5.1 5.2 5.2 5.3 5.3 5.4 5.4 5.5 5.5 5.5 5.6
[127] 5.6 5.6 5.6 5.6 5.6 5.7 5.7 5.7 5.8 5.8 5.8 5.9 5.9 6.0 6.0 6.1 6.1 6.1
[145] 6.3 6.4 6.6 6.7 6.7 6.9
```

```
> stem(plength)
```

The decimal point is at the |

```
1 | 01223333333444444444444444
1 | 555555555555556666666777799
2 |
2 |
3 | 033
3 | 55678999
4 | 000001112222334444
4 | 55555555666777778888999999
```

```

5 | 0000111111111223344
5 | 55566666677788899
6 | 0011134
6 | 6779

```

```
> stem(data1)
```

The decimal point is 2 digit(s) to the left of the |

```

96 | 8
98 | 21389
100 | 44702579
102 | 157799
104 | 227359
106 |
108 | 1
110 | 7
112 | 1

```

```
> stem(data1, scale = 2)
```

The decimal point is 2 digit(s) to the left of the |

```

96 | 8
97 |
98 | 2
99 | 1389
100 | 447
101 | 02579
102 | 1
103 | 57799
104 | 227
105 | 359
106 |
107 |
108 | 1

```

```
109 |  
110 | 7  
111 |  
112 |  
113 | 1
```

```
> stem(data2)
```

The decimal point is at the |

```
2 | 13502469  
4 | 1379  
6 | 1468157  
8 | 0
```

```
> stem(data2, scale = 2)
```

The decimal point is at the |

```
2 | 135  
3 | 02469  
4 | 13  
5 | 79  
6 | 1468  
7 | 157  
8 | 0
```

```
> stem(data2, scale = 4)
```

The decimal point is at the |

```
2 | 13  
2 | 5  
3 | 024  
3 | 69  
4 | 13
```

4		
5		
5		79
6		14
6		68
7		1
7		57
8		0



## 1.4 Boxplots (or Box-and-Whisker Plots)

The R help page for `boxplot` indicates:







“Produce box-and-whisker plot(s) of the given (grouped) values.

`range`: this determines how far the plot whiskers extend out from the box. If `range` is positive, the whiskers extend to the most extreme data point which is no more than `range` times the interquartile range from the box. A value of zero causes the whiskers to extend to the data extremes.”



The default for `range` is 1.5.

?, p. 115, further specify: “A *boxplot* is a way to look at the overall shape of a set of data. The central box shows the data between the ‘hinges’ (roughly quartiles), with the median represented by a line. ‘Whiskers’ go out to the extremes of the data, and very extreme points are shown by themselves.”

```
> par(mfrow = c(2, 3))
> boxplot(plength)
> boxplot(plength, range = 0)
> boxplot(plength, range = 0.1)
> boxplot(plength ~ iris$Species)
> boxplot(plength ~ iris$Species, range = 0)
> boxplot(plength ~ iris$Species, range = 0.5)
```



```
> par(mfrow = c(1, 2))
> boxplot(data1)
> boxplot(data2)
```









## 1.5 Violin Plots

The R help page for `vioplot` indicates:



“Produce violin plot(s) of the given (grouped) values.

A violin plot is a combination of a box plot and a kernel density plot. Specifically, it starts with a box plot. It then adds a rotated kernel density plot to each side of the box plot.”

```
> library(vioplot)
> par(mfrow = c(2, 3))
> vioplot(plength)
> vioplot(plength, range = 0)
> vioplot(plength, range = 0.1)
> vioplot(plength[iris$Species == "setosa"],
+         plength[iris$Species == "versicolor"],
+         plength[iris$Species == "virginica"])
> vioplot(plength[iris$Species == "setosa"],
+         plength[iris$Species == "versicolor"],
+         plength[iris$Species == "virginica"],
+         horizontal = TRUE)
> vioplot(plength[iris$Species == "setosa"],
+         plength[iris$Species == "versicolor"],
+         plength[iris$Species == "virginica"],
+         horizontal = TRUE,
+         names = c("Setosa", "Versicolor", "Virginica"))
```



```
> par(mfrow = c(1, 2))
> vioplot(data1)
> vioplot(data2)
```




## 1.6 Letter Value Boxplots

The R help page for LVboxplot indicates:

“An extension of standard boxplots which draws  $k$  letter statistics. Conventional boxplots (Tukey 1977) are useful displays for conveying rough information about the central 50% of the data and the extent of the data.


For moderate-sized data sets ( $n < 1000$ ), detailed estimates of tail behavior beyond the quartiles may not be trustworthy [...]. Large data sets [...] can be expected to present a large number of ‘outliers’ [...].”

```
> library(lvplot)
> par(mfrow = c(2, 4))
> boxplot(plength, horizontal = TRUE)
> LVboxplot(plength)
> LVboxplot(plength, k = 1)
> LVboxplot(plength, k = 2)
> LVboxplot(plength, k = 3)
> LVboxplot(plength, k = 4)
> LVboxplot(plength, k = 5)
> LVboxplot(plength, k = 6)
```



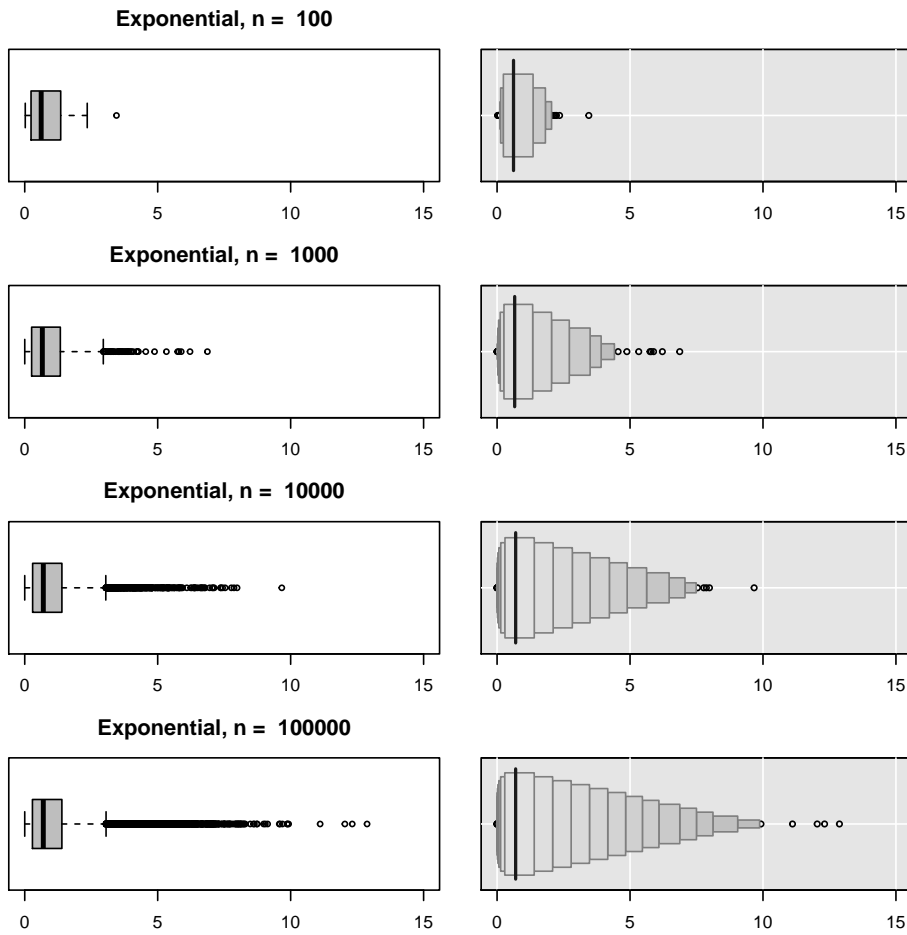
```
> par(mfrow = c(1, 2))
> LVboxplot(data1)
> LVboxplot(data2)
```





Additional example for samples from an exponential distribution, based on LVboxplot help page:

```
> par(mfrow = c(4, 2), mar = c(2, 1, 3, 1))
> for (i in 1:4) {
+   x <- rexp(10^(i + 1))
+   boxplot(x, col = "grey", horizontal = TRUE,
+           ylim = c(0, 15))
+   title(paste("Exponential, n = ", length(x)))
+   LVboxplot(x, col = "grey", xlab = "",
+             xlim = c(0, 15))
+ }
```



## 1.7 Dot Charts for Univariate Data

The R help page for dotchart indicates:

“Draw a Cleveland dot plot.”

The R help page for UScereal(MASS) indicates:

### “Nutritional and Marketing Information on US Cereals:

The UScereal data frame has 65 rows and 11 columns. The data come from the 1993 ASA Statistical Graphics Exposition, and are taken from the mandatory F&DA food label. The data have been normalized here to a portion of one American cup. ”

```
> library(MASS) # for cereal data
> data(UScereal)
> head(UScereal)
```

	mfr	calories	protein	fat	sodium	fibre
100% Bran	N	212.1212	12.121212	3.030303	393.9394	30.303030
All-Bran	K	212.1212	12.121212	3.030303	787.8788	27.272727
All-Bran with Extra Fiber	K	100.0000	8.000000	0.000000	280.0000	28.000000
Apple Cinnamon Cheerios	G	146.6667	2.666667	2.666667	240.0000	2.000000
Apple Jacks	K	110.0000	2.000000	0.000000	125.0000	1.000000
Basic 4	G	173.3333	4.000000	2.666667	280.0000	2.666667
		carbo	sugars	shelf	potassium	vitamins
100% Bran		15.15152	18.18182	3	848.48485	enriched
All-Bran		21.21212	15.15151	3	969.69697	enriched
All-Bran with Extra Fiber		16.00000	0.00000	3	660.00000	enriched
Apple Cinnamon Cheerios		14.00000	13.33333	1	93.33333	enriched
Apple Jacks		11.00000	14.00000	2	30.00000	enriched
Basic 4		24.00000	10.66667	3	133.33333	enriched

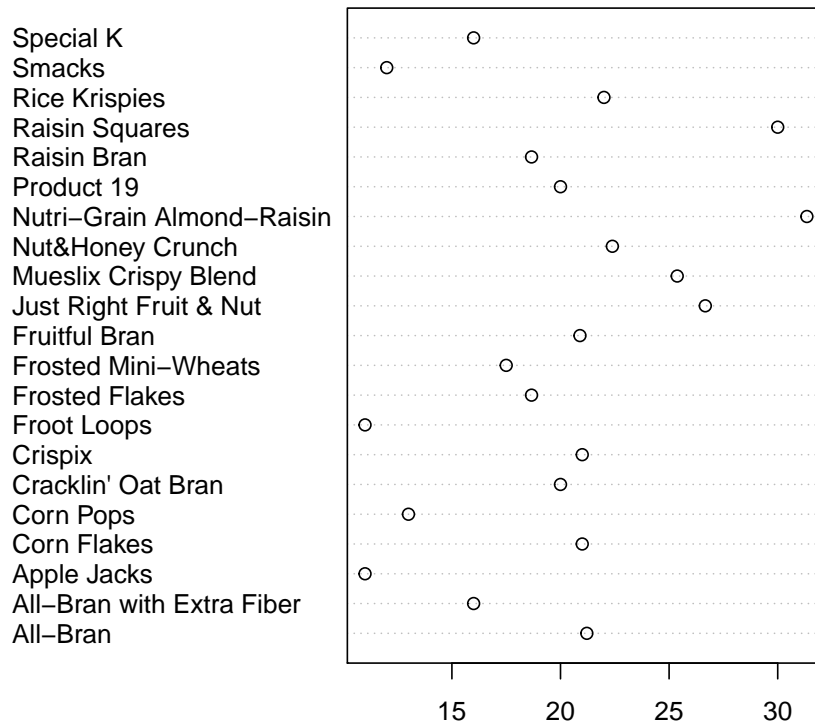
```
> Kel.carbs <- UScereal[UScereal$mfr == "K", 7]
> Kel.carbs
```

```
[1] 21.21212 16.00000 11.00000 21.00000 13.00000 20.00000 21.00000 11.00000
[9] 18.66667 17.50000 20.89552 26.66667 25.37313 22.38806 31.34328 20.00000
[17] 18.66667 30.00000 22.00000 12.00000 16.00000
```


```
> names(Kel.carbs) <- row.names(UScereal[UScereal$mfr == "K", ])
> Kel.carbs
```

All-Bran	All-Bran with Extra Fiber	Apple Jacks
21.21212	16.00000	11.00000
Corn Flakes	Corn Pops	Cracklin' Oat Bran
21.00000	13.00000	20.00000
Crispix	Froot Loops	Frosted Flakes
21.00000	11.00000	18.66667
Frosted Mini-Wheats	Fruitful Bran	Just Right Fruit & Nut
17.50000	20.89552	26.66667
Mueslix Crispy Blend	Nut&Honey Crunch	Nutri-Grain Almond-Raisin
25.37313	22.38806	31.34328
Product 19	Raisin Bran	Raisin Squares
20.00000	18.66667	30.00000
Rice Krispies	Smacks	Special K
22.00000	12.00000	16.00000


```
> dotchart(Kel.carbs)
```



```
> dotchart(sort(Kel.carbs))
```




```
> dotchart(sort(Kel.carbs), xlim = c(10, 35),
+   xlab = "g carbohydrates per 1 cup serving")
```




Now, use the lattice package to produce similar graphs:

```
> library(lattice)
> dotplot(Kel.carbs) # from lattice library
```



```
> dotplot(sort(Kel.carbs), xlim = c(9, 36),
+   xlab = "g carbohydrates per 1 cup serving")
```






Finally, use ggplot2 to produce similar graphs:

```
> library(ggplot2)
> Kel.carbsdf <- data.frame(Brand = as.factor(names(Kel.carbs)), Carbs = Kel.carbs)
> Kel.carbsdf
```

	Brand	Carbs
All-Bran	All-Bran	21.21212
All-Bran with Extra Fiber	All-Bran with Extra Fiber	16.00000
Apple Jacks	Apple Jacks	11.00000
Corn Flakes	Corn Flakes	21.00000
Corn Pops	Corn Pops	13.00000
Cracklin' Oat Bran	Cracklin' Oat Bran	20.00000
Crispix	Crispix	21.00000
Froot Loops	Froot Loops	11.00000

Frosted Flakes	Frosted Flakes	18.66667
Frosted Mini-Wheats	Frosted Mini-Wheats	17.50000
Fruitful Bran	Fruitful Bran	20.89552
Just Right Fruit & Nut	Just Right Fruit & Nut	26.66667
Mueslix Crispy Blend	Mueslix Crispy Blend	25.37313
Nut&Honey Crunch	Nut&Honey Crunch	22.38806
Nutri-Grain Almond-Raisin	Nutri-Grain Almond-Raisin	31.34328
Product 19	Product 19	20.00000
Raisin Bran	Raisin Bran	18.66667
Raisin Squares	Raisin Squares	30.00000
Rice Krispies	Rice Krispies	22.00000
Smacks	Smacks	12.00000
Special K	Special K	16.00000

```
> ggplot(Kel.carbsdf, aes(x = Carbs, y = Brand)) +
+   geom_point()
```




```
> # reorder by carbs
> Kel.carbsdf$Brand <- factor(Kel.carbsdf$Brand,
+                             levels = Kel.carbsdf$Brand[order(Kel.carbsdf$Carbs)])
> Kel.carbsdf$Brand
```

```
[1] All-Bran          All-Bran with Extra Fiber
[3] Apple Jacks       Corn Flakes
[5] Corn Pops         Cracklin' Oat Bran
[7] Crispix           Froot Loops
[9] Frosted Flakes    Frosted Mini-Wheats
[11] Fruitful Bran     Just Right Fruit & Nut
[13] Mueslix Crispy Blend Nut&Honey Crunch
[15] Nutri-Grain Almond-Raisin Product 19
[17] Raisin Bran       Raisin Squares
[19] Rice Krispies     Smacks
```

[21] Special K

21 Levels: Apple Jacks Froot Loops Smacks ... Nutri-Grain Almond-Raisin

```
> ggplot(Kel.carbsdf, aes(x = Carbs, y = Brand)) +  
+   geom_point()
```




The web site at <https://www.r-bloggers.com/revisiting-clevelands-the-elements-of-graphing-d> provides some useful suggestions how to further improve the appearance of these graphs:

```
> # make some further adjustments to the appearance of the grid  
> ggplot(Kel.carbsdf, aes(x = Carbs, y = Brand)) +  
+   geom_point() +  
+   theme(  
+     # remove the vertical grid lines  
+     panel.grid.major.x = element_blank(),
```

```

+   panel.grid.minor.x = element_blank(),
+   # explicitly set the horizontal grid lines
+   panel.grid.major.y = element_line(linetype = 3, color = "darkgray"),
+   axis.text.y = element_text(size = rel(0.8))) +
+   xlab("Carbs (g carbohydrates per 1 cup serving)")
> # interactive version
> library(plotly)
> ggplotly()

```



The R help page for barley in the lattice package indicates:

**“Yield data from a Minnesota barley trial:**


Total yield in bushels per acre for 10 varieties at 6 sites in each of two years.”

```
> library(lattice) # for barley data
> data(barley)
> head(barley)
```

	yield	variety	year	site
1	27.00000	Manchuria	1931	University Farm
2	48.86667	Manchuria	1931	Waseca
3	27.43334	Manchuria	1931	Morris
4	39.93333	Manchuria	1931	Crookston
5	32.96667	Manchuria	1931	Grand Rapids
6	28.96667	Manchuria	1931	Duluth

```
> dotplot(variety ~ yield | site, data = barley, groups = year,
+ key = simpleKey(levels(barley$year), space = "right"),
+ xlab = "Barley Yield (bushels/acre)",
+ aspect = 0.5, layout = c(1, 6), ylab = NULL)
```





**Question:**

What is the most striking (unusual) feature in these plots? Look carefully!



## 1.8 Kernel Density Plots for Univariate Data (with Rug Plot)

The R help page for density indicates:

### “Kernel Density Estimation:

The (S3) generic function `density` computes kernel density estimates. Its default method does so with the given kernel and bandwidth for univariate observations. [...]

`bw`: the smoothing bandwidth to be used. The kernels are scaled such that this is the standard deviation of the smoothing kernel. (Note this differs from the reference books cited below, and from S-PLUS.)

`bw` can also be a character string giving a rule to choose the bandwidth. See `bw.nrd`.

The specified (or computed) value of `bw` is multiplied by `adjust`.”

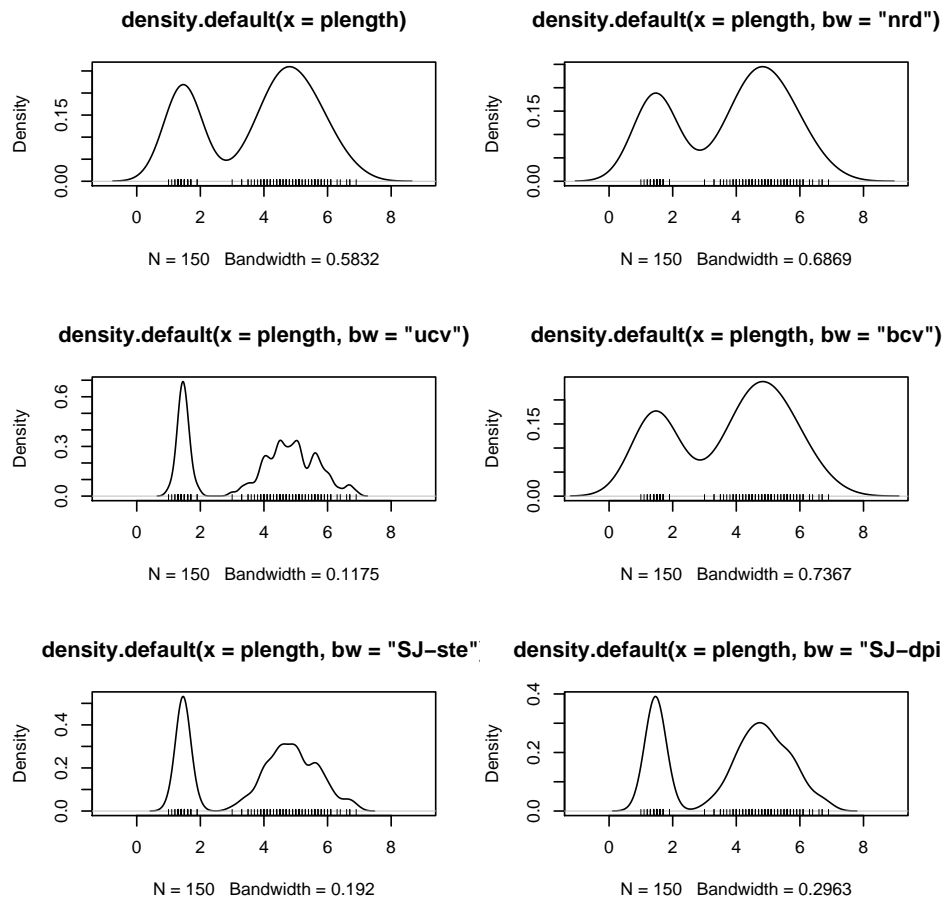
The R help page for `rug` indicates:

“Adds a rug representation (1-d plot) of the data to the plot.”

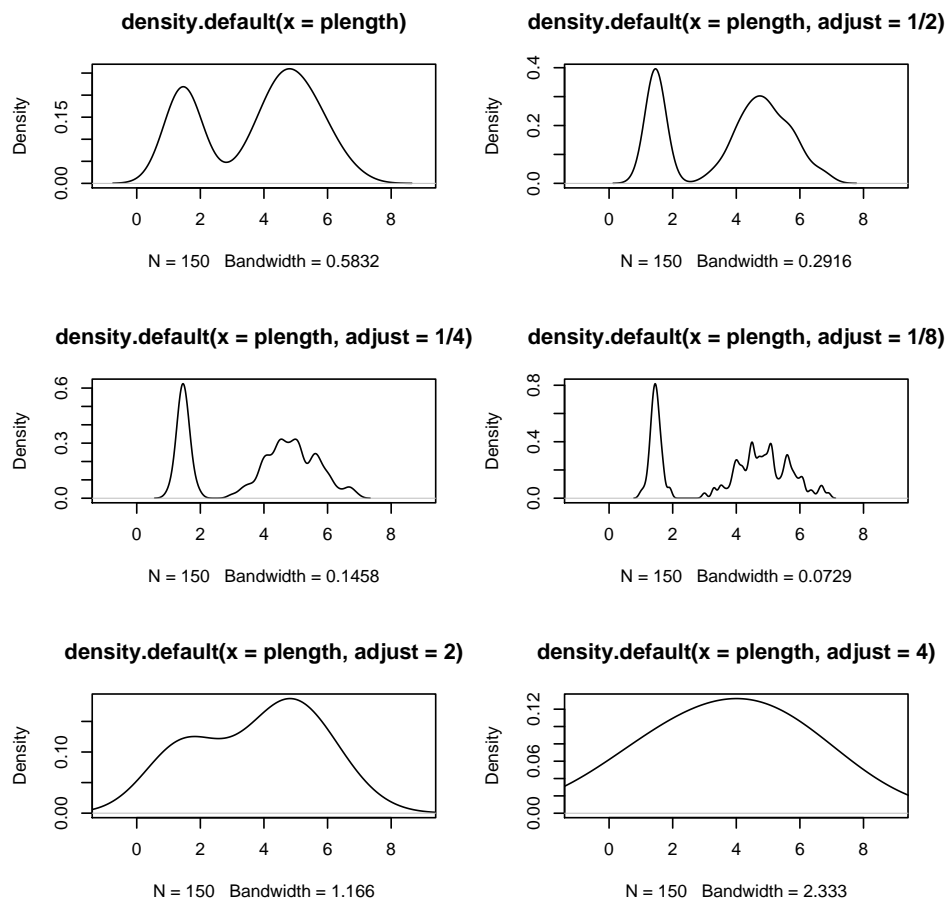
?, p. 548, further specify: *“rug: [...] a univariate histogram or rugplot is displayed along the base of each plot, showing the occurrence of each x-value; ties are broken by jittering.”*

```
> par(mfrow = c(3, 2))
> plot(density(plength), xlim = c(-1, 9))           # default nrd0 bw
> rug(plength, ticksize = 0.05)
> plot(density(plength, bw = "nrd"), xlim = c(-1, 9)) # normal reference rule bw
> rug(plength, ticksize = 0.05)
> plot(density(plength, bw = "ucv"), xlim = c(-1, 9))
>                                                     # unbiased cross-validation rule bw
> rug(plength, ticksize = 0.05)
> plot(density(plength, bw = "bcv"), xlim = c(-1, 9))
>                                                     # biased cross-validation rule bw
> rug(plength, ticksize = 0.05)
> plot(density(plength, bw = "SJ-ste"), xlim = c(-1, 9))
>                                                     # Sheather-Jones ("solve-the-equation") bw
> rug(plength, ticksize = 0.05)
```

```
> plot(density(plength, bw = "SJ-dpi"), xlim = c(-1, 9))
>                                     # Sheather-Jones ("direct plug-in") bw
> rug(plength, ticksize = 0.05)
```

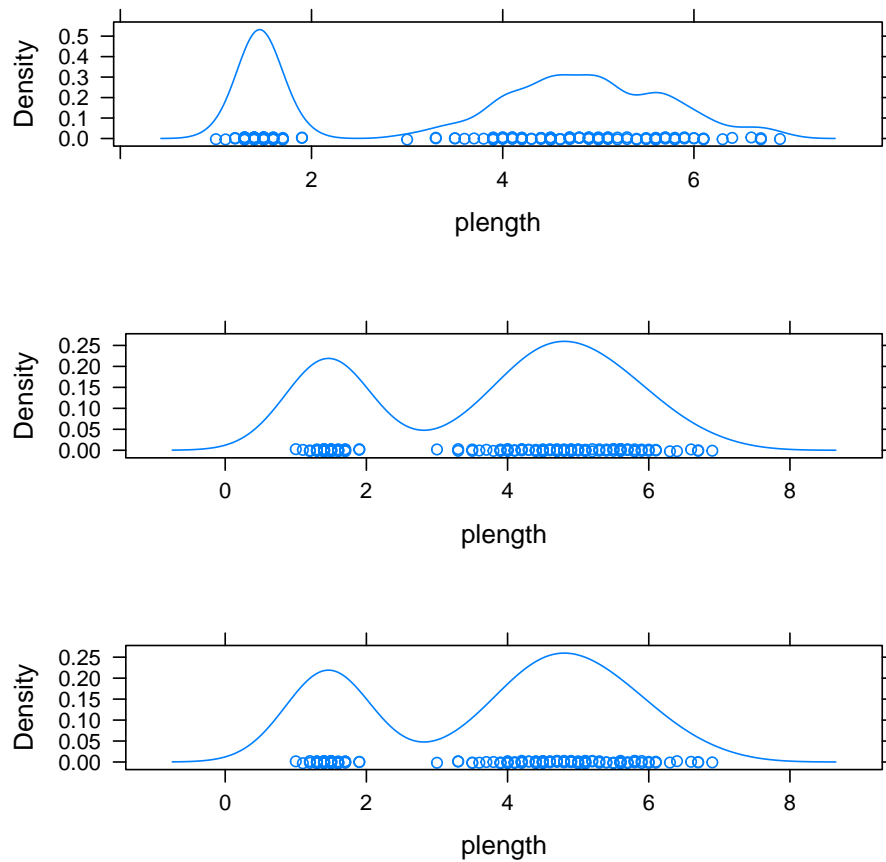


```
> par(mfrow = c(3, 2))
> plot(density(plength), xlim = c(-1, 9))
> plot(density(plength, adjust = 1/2), xlim = c(-1, 9)) #adjust default bandwidth
> plot(density(plength, adjust = 1/4), xlim = c(-1, 9))
> plot(density(plength, adjust = 1/8), xlim = c(-1, 9))
> plot(density(plength, adjust = 2), xlim = c(-1, 9))
> plot(density(plength, adjust = 4), xlim = c(-1, 9))
```



Now, use the lattice package and produce similar graphics:

```
> library(lattice)
> dp1 <- densityplot(plength) # lattice, takes all parameters from density (above)
> dp2 <- densityplot(plength, n = 512)
> dp3 <- densityplot(plength, n = 512, bw = "SJ")
> # arrange the three plots vertically
> print(dp1, position = c(0, 0, 1, 0.33), more = TRUE)
> print(dp2, position = c(0, 0.33, 1, 0.66), more = TRUE)
> print(dp3, position = c(0, 0.66, 1, 0.99))
```



In *densityplot*, *n* is the “number of points at which density is to be evaluated” and the default is 50.

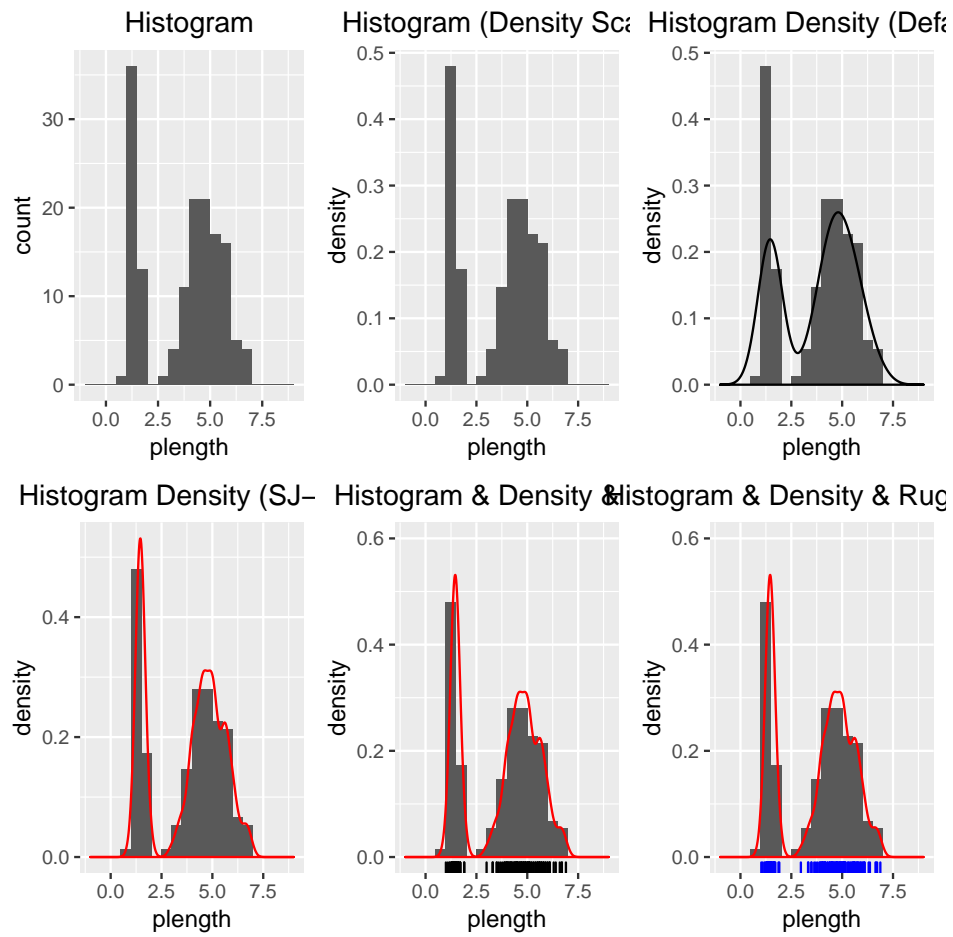
Now in *ggplot2*. The web page <https://stackoverflow.com/questions/32468772/histogram-with-a-jittery-rug> provided some suggestions how to add the jittered rug.

```
> library(ggplot2)
> library(gridExtra)
> h1 <- ggplot(iris, aes(x = plength)) +
+   xlim(-1, 9) +
+   geom_histogram(binwidth = 0.5, center = 0.25) +
+   ggtitle("Histogram") +
+   theme(plot.title = element_text(hjust = 0.5))
> h2 <- ggplot(iris, aes(x = plength)) +
+   xlim(-1, 9) +
```

```

+   geom_histogram(aes(y = ..density..), binwidth = 0.5, center = 0.25) +
+   ggtitle("Histogram (Density Scale)") +
+   theme(plot.title = element_text(hjust = 0.5))
> h3 <- ggplot(iris, aes(x = plength)) +
+   xlim(-1, 9) +
+   geom_histogram(aes(y = ..density..), binwidth = 0.5, center = 0.25) +
+   ggtitle("Histogram Density (Default)") +
+   theme(plot.title = element_text(hjust = 0.5)) +
+   geom_density()
> h4 <- ggplot(iris, aes(x = plength)) +
+   xlim(-1, 9) +
+   geom_histogram(aes(y = ..density..), binwidth = 0.5, center = 0.25) +
+   ggtitle("Histogram Density (SJ-ste)") +
+   theme(plot.title = element_text(hjust = 0.5)) +
+   geom_density(bw = "SJ-ste", col = "red")
> h5 <- ggplot(iris, aes(x = plength)) +
+   xlim(-1, 9) +
+   ylim(0, 0.6) +
+   geom_histogram(aes(y = ..density..), binwidth = 0.5, center = 0.25) +
+   ggtitle("Histogram & Density & Rug") +
+   theme(plot.title = element_text(hjust = 0.5)) +
+   geom_density(bw = "SJ-ste", col = "red") +
+   geom_rug()
> h6 <- ggplot(iris, aes(x = plength)) +
+   xlim(-1, 9) +
+   ylim(0, 0.6) +
+   geom_histogram(aes(y = ..density..), binwidth = 0.5, center = 0.25) +
+   ggtitle("Histogram & Density & Rug (Final)") +
+   theme(plot.title = element_text(hjust = 0.5)) +
+   geom_density(bw = "SJ-ste", col = "red") +
+   geom_rug(sides = "b", aes(y = 0.4), position = "jitter", col = "blue")
> grid.arrange(h1, h2, h3, h4, h5, h6, nrow = 2)

```



## 1.9 Quantile–Quantile Plots (Q–Q Plots)

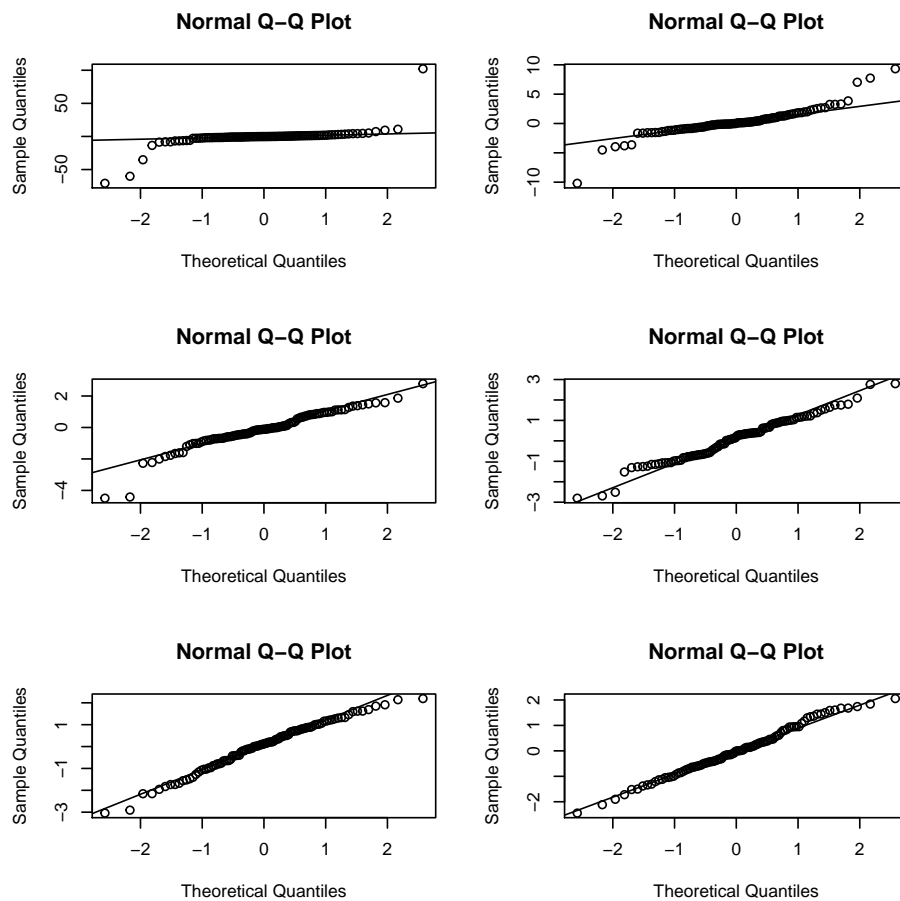
One of the best ways to compare the distribution of a sample  $\underline{x}$  of size  $n$  with an assumed theoretical distribution  $F$  is to use a Quantile–Quantile Plot (Q–Q Plot). In such a plot, we plot the pairs of points

$$\left( F^{-1} \left( \frac{i - 0.5}{n} \right), x_{(i)} \right), \quad i = 1, \dots, n.$$

### Example 1:

Convergence of a  $t$  distribution with  $df$  degrees of freedom towards a normal distribution.

```
> # set seed of random number generator to be able to reproduce results
> set.seed(2)
> par(mfrow = c(3, 2))
> tdf1 <- rt(100, df = 1)
> qqnorm(tdf1)
> qqline(tdf1)
> tdf2 <- rt(100, df = 2)
> qqnorm(tdf2)
> qqline(tdf2)
> tdf5 <- rt(100, df = 5)
> qqnorm(tdf5)
> qqline(tdf5)
> tdf10 <- rt(100, df = 10)
> qqnorm(tdf10)
> qqline(tdf10)
> tdf20 <- rt(100, df = 20)
> qqnorm(tdf20)
> qqline(tdf20)
> tdf30 <- rt(100, df = 30)
> qqnorm(tdf30)
> qqline(tdf30)
```



### Note:

The closer the points from the sample fall to a straight line, the closer the sample distribution and the theoretical distribution are related. However, here, the greater spread of the extreme quantiles for the sample (for  $df = 1, 2, 5, 10, 20$ ) is an indicator of a long-tailed distribution.

A quick glance at Q-Q-plots via ggplot2, with additional suggestions from [http://ggplot2.tidyverse.org/reference/geom\\_qq.html](http://ggplot2.tidyverse.org/reference/geom_qq.html) and <https://gist.github.com/slowkow/9041570>.


```
> library(ggplot2)
> library(gridExtra)
> tdf1df <- data.frame(y = tdf1)
> q1 <- ggplot(tdf1df, aes(sample = y)) +
+   stat_qq() +
+   geom_abline(intercept = 0, slope = 1, alpha = 0.5)
```



```

> tdf30df <- data.frame(y = tdf30)
> q30 <- ggplot(tdf30df, aes(sample = y)) +
+   stat_qq() +
+   geom_abline(intercept = 0, slope = 1, alpha = 0.5)
> grid.arrange(q1, q30, nrow = 1)

```



### Example 2:

Recall: What is the relationship between exponential distributions and Gamma distributions? Verify this graphically!

```

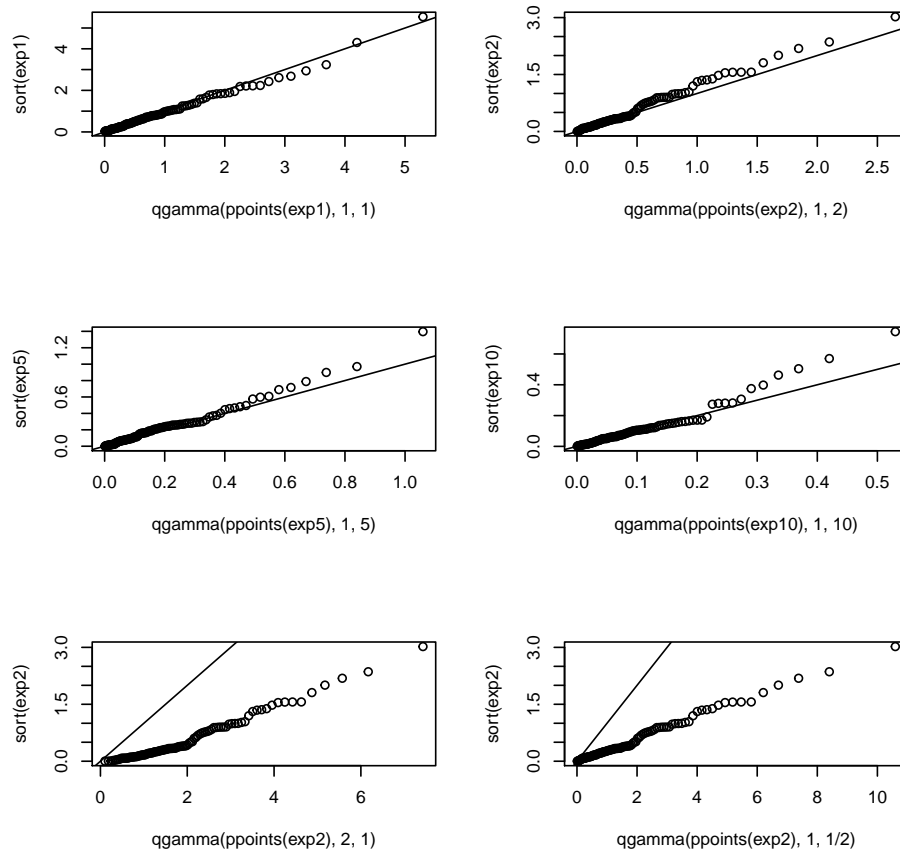
> par(mfrow = c(3, 2))
> set.seed(4)
> exp1 <- rexp(100, rate = 1)
> plot(qgamma(ppoints(exp1), 1, 1), sort(exp1))
> abline(0, 1)

```

```

> exp2 <- rexp(100, rate = 2)
> plot(qgamma(ppoints(exp2), 1, 2), sort(exp2))
> abline(0, 1)
> exp5 <- rexp(100, rate = 5)
> plot(qgamma(ppoints(exp5), 1, 5), sort(exp5))
> abline(0, 1)
> exp10 <- rexp(100, rate = 10)
> plot(qgamma(ppoints(exp10), 1, 10), sort(exp10))
> abline(0, 1)
> # Some major misspecifications
>
> # a) Swapping shape and rate parameters
> plot(qgamma(ppoints(exp2), 2, 1), sort(exp2))
> abline(0, 1)
> # b) Using 1/rate instead
> plot(qgamma(ppoints(exp2), 1, 1/2), sort(exp2))
> abline(0, 1)

```




### Example 3:


Compare iris plenth sample data with an underlying assumed normal distribution. For which of the species is the assumption of normality justified?

```
> par(mfrow = c(2, 2))
> qqnorm(plenth)
> qqnorm(plenth[1:50])
> qqnorm(plenth[51:100])
> qqnorm(plenth[101:150])
```


**Normal Q-Q Plot**




**Normal Q-Q Plot**



**Normal Q-Q Plot**



**Normal Q-Q Plot**



## 1.10 Empirical Cumulative Distribution Functions (ECDFs)

Recall from Stat 6720:

Definition 7.1.3:

Let  $X_1, \dots, X_n$  be a sample of size  $n$  from a population with distribution  $F$ . The function

$$\hat{F}_n(x) = \frac{1}{n} \sum_{i=1}^n I_{(-\infty, x]}(X_i)$$

is called **empirical cumulative distribution function** (empirical cdf, ECDF). ■

Theorem 7.1.7: Glivenko–Cantelli Theorem

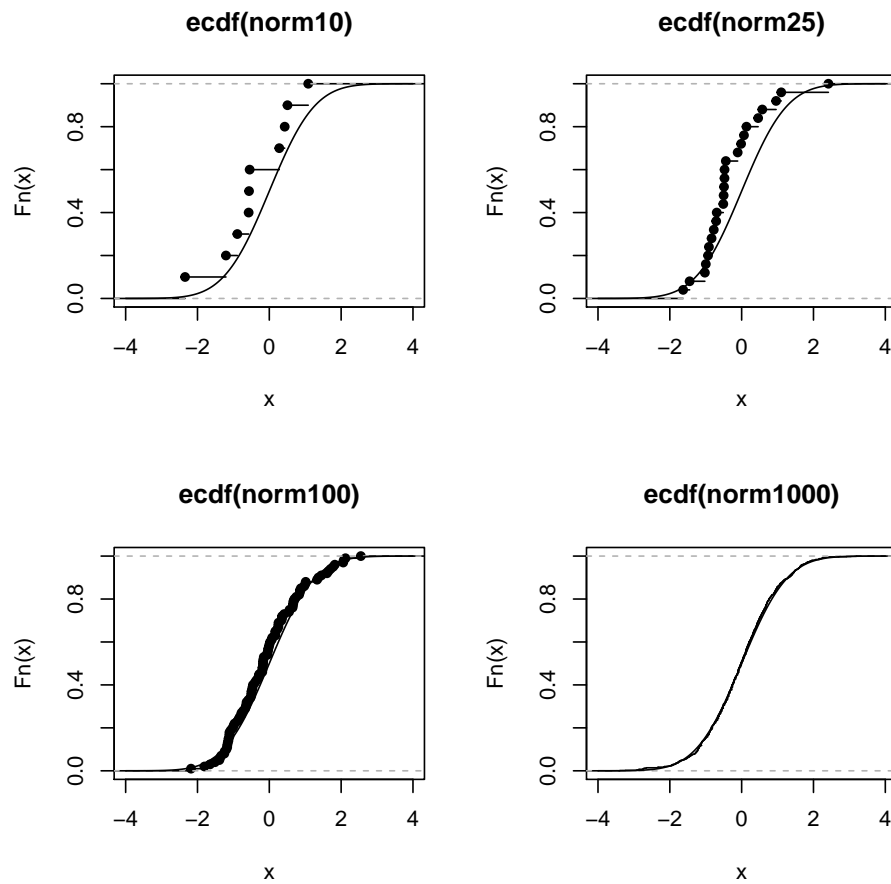
$\hat{F}_n(x)$  converges uniformly to  $F(x)$ , i.e., it holds for all  $\epsilon > 0$  that

$$\lim_{n \rightarrow \infty} P\left(\sup_{-\infty < x < \infty} |\hat{F}_n(x) - F(x)| > \epsilon\right) = 0.$$

■

Verify this theorem for samples from a normal distribution:

```
> par(mfrow = c(2, 2))
> set.seed(1234)
> xvals <- seq(-4, 4, 0.01)
> norm10 <- rnorm(10)
> plot(ecdf(norm10), xlim = c(-4, 4))
> lines(xvals, pnorm(xvals))
> norm25 <- rnorm(25)
> plot(ecdf(norm25), xlim = c(-4, 4))
> lines(xvals, pnorm(xvals))
> norm100 <- rnorm(100)
> plot(ecdf(norm100), xlim = c(-4, 4))
> lines(xvals, pnorm(xvals))
> norm1000 <- rnorm(1000)
> plot(ecdf(norm1000), xlim = c(-4, 4))
> lines(xvals, pnorm(xvals))
```



```


> # Animation
> par(mfrow = c(1, 1))
> xvals <- seq(-4, 4, 0.01)
> pnormxvals <- pnorm(xvals)
> normgrow <- NULL
> for (i in 1:150) {
+   normgrow <- c(normgrow, rnorm(1))
+   plot(ecdf(normgrow), xlim = c(-4, 4), main = length(normgrow))
+   lines(xvals, pnormxvals, col = "red")
+   Sys.sleep(0.1)
+ }

```

### Example:

And here the ECDF for iris plength.

```
> plot(ecdf(plength))
```



## 1.11 Graphics and Small Sample Sizes

### Worksheet

Your Name: \_\_\_\_\_

Question:

The data shown in these four histograms originate from which distribution?

The corresponding distributions are:

Upper left: \_\_\_\_\_

Upper right: \_\_\_\_\_

Lower left: \_\_\_\_\_

Lower right: \_\_\_\_\_



# Worksheet

**Your Name:** \_\_\_\_\_

**Question:**

Do the data shown in these four qqplots follow a normal distribution?

So, does a particular qqplot suggest that the data originate from a normal distribution?

Circle your answer:

Upper left:    **yes**    /    **no**

Upper right:    **yes**    /    **no**

Lower left:    **yes**    /    **no**

Lower right:    **yes**    /    **no**

### Answers:

The following R code was used to create the figures shown on the previous two pages:

```
> set.seed(1234)
> jpeg("Chapter5_unknown_hist.jpg")
> par(mfrow = c(2, 2))
> xvect1 <- NULL
> for (i in 1:4) {
+   x <- rnorm(10)
+   xvect1 <- c(xvect1, x)
+   hist(x, main = "Unknown Distribution",
+       xlab = "x-values")
+ }
> dev.off()
```

pdf

2


```
> set.seed(1234)
> jpeg("Chapter5_unknown_qqplot.jpg")
> par(mfrow = c(2, 2))
> xvect2 <- NULL
> for (i in 1:4) {
+   x <- rnorm(10)
+   xvect2 <- c(xvect2, x)
+   qqnorm(x, main = "Unknown Distribution",
+       xlab = "Normal Quantiles")
+ }
> dev.off()
```

pdf


2

When we jointly plot all 40 observations, we start to see that the underlying distribution indeed is a normal distribution. In fact, all eight plots on the previous two pages show 10 samples each drawn from the standard normal distribution!

```
> # Plot all data combined  
>  
> par(mfrow = c(1, 1))  
> hist(xvect1)
```



```
> par(mfrow = c(1, 1))  
> qqnorm(xvect2)
```



## 1.12 Further Reading

Additional sources for Trellis Graphics are:

- Trellis displays at Bell Labs (in S/S-Plus; needs some updates for modern R): <http://ect.bell-labs.com/sl/project/trellis/>
- ?, Chapter 4: <https://www.stat.auckland.ac.nz/~paul/RGraphics/chapter4.pdf>
- William G. Jacoby's web page on dot plots: <http://polisci.msu.edu/jacoby/research/statgrph/tpm/sdlist.html>

— THE END —